

Exercise 8: Implementing the PUCT Algorithm

January 16, 2025

In this project, you will implement the PUCT algorithm. This exercise will build on Exercises 4 and 7.

1. Adapting MCTS for Your Board Game

- Adapt your MCTS code from Exercise 4 to your chosen board game from Exercise 7.
- Debug thoroughly to ensure it plays the game correctly and performs well.

2. Creating the Neural Network Class

- Write a Python class called `GameNetwork` for your neural network.
- The network must have a value head estimating win probability and a policy head predicting action probabilities.
- Ensure input and output dimensions match `encode` and `decode` from Exercise 7.
- Include methods for saving and loading network weights.

3. Implementing PUCT Algorithm

- Write `PUCTNode` and `PUCTPlayer` classes by adapting `MCTSNode` and `MCTSPlayer`.
- Key changes:
 - Replace random rollouts with evaluations from the neural network.
 - Each node maintains a prior probability P from the policy head of its parent.
 - Each node maintains Q (average value) and N (visit count). Avoid calling it wins.
 - Use the PUCT scoring formula:

$$U(s, a) = Q(s, a) + c_{puct} \cdot P(s, a) \cdot \frac{\sqrt{N(s)}}{1 + N(s, a)}$$

4. **Pre-training the Neural Network**

- Use your MCTSPlayer to generate 10,000 self-play games without network guidance.
- Train the value head using game outcomes.
- Train the policy head using visit counts at the root node.